





# Image Analysis & Computer Vision

## Visual analysis in sport events: Tennis

2<sup>nd</sup> Semester of 2023/2024 A.Y.

Mihai-Viorel Grecu  
Paolo Riva  
Michelangelo Stasi

# Visual Analysis in sports: Why the interest?



Improved performance through statistics analysis



Significant positive economic impact



Higher quality communication approach of sports analysis



Improved understanding of the sport for viewers



# Visual Analysis in Tennis:

## Motivation?



4<sup>th</sup> most popular sport: 1 billion fans



Staggering increase in popularity over 1 year: 43% - 49%





# Visual Analysis in Tennis: Goal of the project



Trajectory computation of a tennis player through a monocular video



Ball trajectory computation and statistics



# Visual Analysis in Tennis:

## Goal of the project



Trajectory computation of a tennis player through a monocular video

1. Homography identification from the field to the image
2. Human Pose estimation implementation
3. Dynamics of the feet of the tennis players identification
4. Step points collection



State of the art

# State of the art

Homography + Tennis player trajectory tracking + Ball tracking

HOMOGRAPHY IDENTIFICATION  
AND COMPUTATION

OPENCV LIBRARY



HUMAN POSE ESTIMATION  
OPEN POSE



HUMAN DETECTION

+

HUMAN POSE ESTIMATION  
MEDIPIPE POSE



TENNIS BALL TRACKING  
YOU ONLY LOOK ONCE



TRACKNET  
TRACE



# State of the art



## Homography

### HOMOGRAPHY IDENTIFICATION AND COMPUTATION

#### OPENCV LIBRARY



$$s \begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = H \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$



$X'$  = projection of  $X$  in a homogenous coordinate plane  
→ same information, but in transformed perspective

$$\sum_i \left( x'_i - \frac{h_{11}x_i + h_{12}y_i + h_{13}}{h_{31}x_i + h_{32}y_i + h_{33}} \right)^2 + \left( y'_i - \frac{h_{21}x_i + h_{22}y_i + h_{23}}{h_{31}x_i + h_{32}y_i + h_{33}} \right)^2$$

Back projection error minimized through:

**cv2FindHomography** function ⇒

finds  $H$  between the source and destination plane

# State of the art

Tennis player trajectory tracking

## HUMAN POSE ESTIMATION OPEN POSE



- Pulls features from the entire frame
- Confidence maps for each of the 27 body parts of the human pose skeleton
- Body part association
- Human pose skeleton assemble

# State of the art

Tennis player trajectory tracking

HUMAN POSE ESTIMATION

OPEN POSE



HUMAN DETECTION

+

HUMAN POSE ESTIMATION

MEDIAPIPE POSE



- Bounding box computation for human
- 32 Landmarks prediction (single shot approach) + linking
- Pose refinement
- Temporal filtering to smooth out jitter or noise
- Mapping of the keypoints

# State of the art

Ball trajectory tracking

## TENNIS BALL TRACKING YOU ONLY LOOK ONCE



- Bounding boxes prediction
- Class probabilities computation
- Use of a threshold to keep the highest confidence box



In a single  
shot

# State of the art

Ball trajectory tracking

TENNIS BALL TRACKING  
YOU ONLY LOOK ONCE



TRACE +  
TRACKNET

- Probability-like detection heatmap for object tracking
- Upsampling to recover the information loss = pixel-wise prediction
- Use of a threshold to keep the highest confidence box



# Our Implementation

# Our Implementation

---


Homography Computation

# Automatic Homography Computation

## Hough Transform and Probability Hough Transform

First approach :

- Selecting points manually
- Map selected points with real field vertices



# Automatic Homography Computation


## Hough Transform and Probability Hough Transform

Second approach :


- Selecting points ~~manually~~ automatically
- Map selected points with real field vertices



Probability Hough Transform



Dilating lines:  
The lines with most intersections  
get a fill mask command




Eroding Lines

# Automatic Homography Computation

## Hough Transform and Probability Hough Transform

Second approach :

- Selecting points ~~manually~~ automatically
- Map selected points with real field vertices



Hough Transform:  
Threshold – 300 px



Find intersections



# Our Implementation

---


Human Detection & Human Pose Estimation

# Our implementation

Human detection & Human Pose Estimation

HUMAN POSE ESTIMATION  
**OPEN POSE**

HUMAN DETECTION  
+  
HUMAN POSE ESTIMATION  
**MEDIAPIPE POSE**



ACCURACY



TIME



# Human Pose Estimation - OpenPose

Implementation of the human pose estimation method in a straight-forward, out-of-the-box ready approach.

On the other hand, in terms of computational time and noise seen on the output, the model was performing poorly especially in high noise – high movement scenarios, which made the implementation not the most suitable for our case.



# Human Pose Estimation - MediaPipe

Built-in model more accurate and faster than OpenPose model, even if more complex.

Two problems :

- The multi-object detection is possible but really slow if treated in a standard way;
- For efficiency, the estimation has to be done on specific regions of the frame




# Human Detection - MediaPipe

Less effort demanding than the Human Pose Estimation. Pre-trained model implemented through MediaPipe library called *EfficientDet\_Lite0*. Who is the player at the top of the frame and who is the one at the bottom ?

- Discriminant given by the center of the frame wrt to its height

Noise : storing the previous detection, which is updated when the distance with the new detection is below a certain threshold. Regions on which to compute Human Pose Estimation given by the extreme points reached. Each player has its own area.





# Human Pose Estimation - MediaPipe

To determine whether a player was moving or not, we check if the two feet are static or not.

At each frame for each foot we check if the previous **rectified** position is at a distance greater than a certain threshold; if it is, then the foot is moving, otherwise it is static.

- Previous position updated every 5 frames
- Threshold of 5 pixels



# Our Implementation

---


Ball Detection

# Ball Detection: YOLOv8



## Early-stage model testing

- Most used solution for object detection
- Many versions available in terms of suitability
  - No Hardware or Software restrictions → *yolov8x.pt*
- Selection of the object through label filtering → *sports\_ball*



Performance Test required  
in our scenario



# Ball trajectory detection


## YoloV8x Performance Test



0: 384x640 7 persons, 1 chair, 1690.3ms  
Speed: 1.0ms preprocess, 1690.3ms inference, 2.0ms postprocess per image  
Total Frames: 382  
Frames where a tennis ball was detected: 68  
Percentage of frames where a tennis ball was detected: 17.80%

# Ball trajectory detection

TrackNet Model implementation



- ~ 80% accuracy
- Multi-frame detection
- Trained specifically for our scenario
- Trajectory detection



# Ball trajectory detection

Tennis Ball detection

TRACE interface



```
class BallDetector:
    def __init__(self, save_state, out_channels=2):
    def detect_ball(self, frame):
        # Save frame dimensions
        if self.video_width is None:
            self.video_width = frame.shape[1]
            self.video_height = frame.shape[0]
        self.last_frame = self.before_last_frame
        self.before_last_frame = self.current_frame
        self.current_frame = frame.copy()

        # detect only in 3 frames were played
        if self.last_frame is not None:
            # combine the frames into 1 input tensor
            frames = combine_three_frames(self.current_frame, self.before_last_frame, self.last_frame,
                                         self.model_input_width, self.model_input_height)
            frames = (torch.from_numpy(frames) / 255).to(self.device)
            # Inference (forward pass)
            x, y = self.detector.inference(frames)
            if x is not None:
                # Rescale the indices to fit frame dimensions
                x = int(x * (self.video_width / self.model_input_width))
                y = int(y * (self.video_height / self.model_input_height))

                # Check distance from previous location and remove outliers
                if self.xy_coordinates[-1][0] is not None:
                    if np.linalg.norm(np.array([x, y]) - self.xy_coordinates[-1]) > self.threshold_dist:
                        x, y = None, None
                self.xy_coordinates = np.append(self.xy_coordinates, np.array([x, y]), axis=0)
```

# Ball trajectory estimation

## Frame Computation

### Scene from TrackNet



Computed by *processBallTrajectory*:

- Default 0.5 threshold on heatmap
- Error detection:
  - In-sequence (200px thr.)
  - Out-of-sequence (100px thr.)
- Marking of points to-interpolate


# Ball trajectory estimation

## Interpolation

### Interpolation Domain



- Existing ball positions
- Interpolation results



### Interpolate missing values:

- Spline Interpolation
- Dynamic window to use known ball positions around players
- Spline realistically represents ball behaviour here

# Ball trajectory estimation

## Interpolation Result



# Our Implementation



---

Racket Hit Detection




# Racket Hit Detection

Filtering Ball Trajectory: Savitzky – Golay Filter



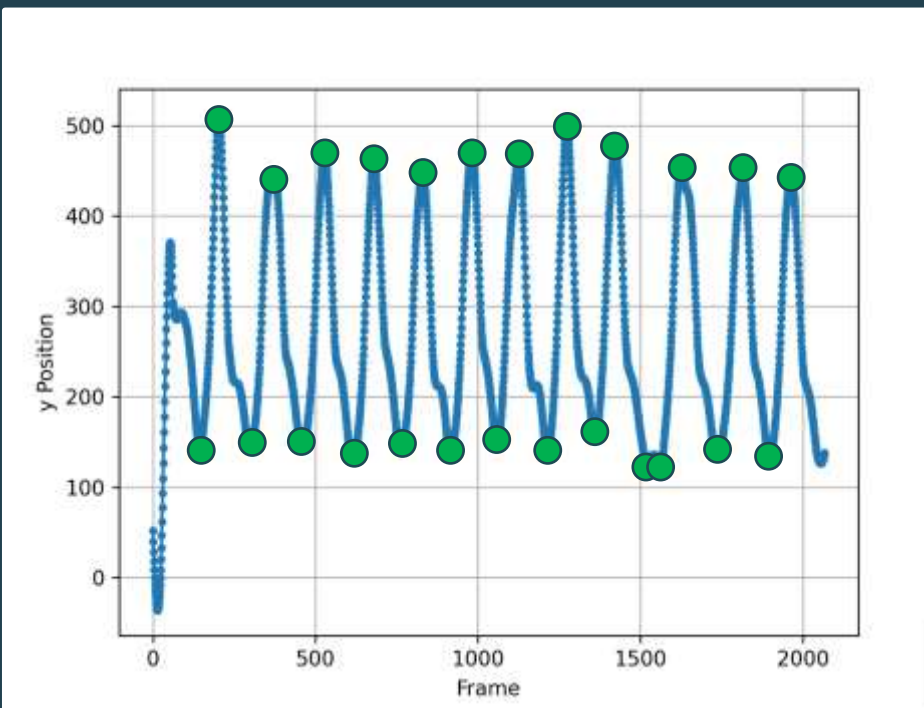
Unfiltered Position



Filtered Result

# Racket Hit Detection

## Filtering Ball Trajectory: Savitzky – Golay Filter



● Detected gradient sign changes



### Outer loop

```
for i in velocity_changes:
    append_flag = False
    min_player_distance = float('inf')
    max_wrist_velocity_sum = 0
    frame_ball_closest_to_player = i

    for j in range(max(0, i - window_around_shot), min(len(ball_positions_array), i + window_around_shot)):
        radiuses = [height_values_top[j] * radius_multiplier, height_values_bot[j] * radius_multiplier]
        if j >= len(ball_positions_array):
            break
        ball_pos = ball_positions_array[j]
        if np.array_equal(ball_pos, np.array([0, 0])):
            break
```

### detect\_racket\_hits

```
if (dist_right_top < max(radiuses[0], default_minimum_radius) or
    dist_left_top < max(radiuses[0], default_minimum_radius) or
    dist_right_bot < max(radiuses[1], default_minimum_radius) or
    dist_left_bot < max(radiuses[1], default_minimum_radius)):

    if j > 0:
        wrist_velocity_right_top = np.linalg.norm(np.array(rightwrist_positions_top[j]) - np.array(rightwrist_positions_top[j - 1]))
        wrist_velocity_left_top = np.linalg.norm(np.array(leftwrist_positions_top[j]) - np.array(leftwrist_positions_top[j - 1]))
        wrist_velocity_right_bot = np.linalg.norm(np.array(rightwrist_positions_bot[j]) - np.array(rightwrist_positions_bot[j - 1]))
        wrist_velocity_left_bot = np.linalg.norm(np.array(leftwrist_positions_bot[j]) - np.array(leftwrist_positions_bot[j - 1]))

        wrist_velocity_sum = wrist_velocity_right_top + wrist_velocity_left_top + wrist_velocity_right_bot + wrist_velocity_left_bot

        if wrist_velocity_sum > max_wrist_velocity_sum:
            max_wrist_velocity_sum = wrist_velocity_sum
            frame_ball_closest_to_player = j

    append_flag = True
```

### Inner Loop



# Our Implementation

---

Statistics and Performance

# Statistics – Distance Travelled & Player's Trajectory

Computed starting from the moving feet detection :

- List of static centers
- Distance as the length of the graph containing all the static centers (in pixels)
- Conversion in meters using ratio computed when computing homography

```
if len(stationary_points_bot)!=0 :
    dist_bot = 0
    cv2.circle(rectified_image, stationary_points_bot[0], 2, (125, 125, 125), cv2.FILLED)
    for z in range(1,len(stationary_points_bot)):
        cv2.circle(rectified_image, stationary_points_bot[z], 2, (125, 125, 125),cv2.FILLED)
        cv2.line(rectified_image, stationary_points_bot[z-1],stationary_points_bot[z], (125,125,125), 3)
        dist_bot += np.linalg.norm(np.array(stationary_points_bot[z]) - np.array(stationary_points_bot[z-1]))/ratiopxpermtr
    dist_bot = np.trunc(dist_bot)
    cv2.putText(frame, "Bottom Player Distance : " + str(dist_bot)+" m", (50,80), cv2.FONT_HERSHEY_SIMPLEX,0.5,(70,150,255), 1)

if len(stationary_points_top)!=0 :
    dist_top = 0
    cv2.circle(rectified_image, stationary_points_top[0], 2, (125, 125, 125), cv2.FILLED)
    for z in range(1,len(stationary_points_top)):
        cv2.circle(rectified_image, stationary_points_top[z], 2, (125, 125, 125), cv2.FILLED)
        cv2.line(rectified_image, stationary_points_top[z-1],stationary_points_top[z], (125,125,125), 3)
        dist_top += np.linalg.norm(np.array(stationary_points_top[z]) - np.array(stationary_points_top[z-1]))/ratiopxpermtr
    dist_top= np.trunc(dist_top)
```

# Statistics – Distance Travelled & Player's Trajectory




*Trajectory Detected*



*Distance covered*

# Statistics

## Average Ball Speed between hits



```
# Pixel distance between the real centers in the transformed perspective (real)
distance = np.linalg.norm(center_start - center_end)
print(f"center_start: {center_start}, center_end: {center_end}")
print(f"Distance: {distance}")

# Time difference in seconds
time_difference = (hit_end - hit_start) / frame_rate
print(f"time_difference: {time_difference}")

# Evaluate average speed of exchange
if time_difference != 0:
    speed = distance / time_difference
else:
    speed = 0
print(f"Speed: {speed}\n")

speeds.append(speed)
```

compute avg ballspeed

Computation between **starting hit**  
and **receiving hit**

# Result analysis

---


Performance and Tests

# Result analysis


## Tests



Australian Open




Wimbledon



Roland Garros

# Result analysis

## Performance




Warping issue due to external factors affecting the homography identification



# Statistics

## Computational Performance



Total processing time: 296.09 s  
Average processing time: 775.10 ms/frame  
Max: 1042.95 ms - Min: 31.03 ms

### Hardware Platform:

- CPU: AMD Ryzen 7 3700X 8-core
- GPU: AMD RX 470 4GB
- RAM: 16GB DDR4 3000MHz CL18
- ROM: 512GB M.2 NVMe SSD

### Software Environment:

- *Python version: 3.12.4 64bit*
- *OS: Windows 11 64bit*

# Conclusions

---

Summary

# Conclusions

## Factors adding difficulty & main challenges

### MAIN CHALLENGES

- Dynamical selection of the frame areas where to estimate the human pose
- Automatic detection of movement associated to players' feet
- Ball Trajectory interpolation and Filtering
- Racket Hit detection through ball trajectory
- Automatic computation of the homography from field to image



- Low quality camera
- Motion blurring



- Low resolution
- Low framerates



- Referees in the background



- Environmental changes: shade

**Factors adding difficulty to the trajectory tracking of the tennis players**